

INDEX

S.NO	TITLE	PAGE NO.
1	Elasticsearch	2
2	Inverted Index	2
3	How data is stored inside elasticsearch?	3
4	Indices API	5
5	Mapping	7
6	Query DSL	9
7	Aggregations	17
8	Index Aliases	22

Elasticsearch Documentation

1. Elasticsearch

- Elasticsearch is a [document](#) oriented search engine.
- It is built on a search software 'Lucene'.
- Underlying data structure in Lucene is known as **Inverted Index**.

2. Inverted Index

- Maps words to the actual document locations of where they occur.
- eg:

Doc 1 : A brown fox jumps.

Doc 2 : The fox quickly jumps over dog.

<u>Tokens/Terms</u>	<u>Doc 1</u>	<u>Doc 2</u>
A	X	
brown	X	
fox	X	X
jumps	X	X
The		X
quickly		X
over		X
dog		X

3. How data is stored inside elasticsearch?

3.1. Document

- It is smallest unit of data that can be indexed(stored) in elasticsearch.
- It is expressed as a JSON file.
- A JSON file consists of Key(field)-Value pairs.
- Structure of JSON file :-
eg:

```
{
  "total" : 2,
  "Employees" : [
    {
      "name" : "David",
      "id" : 1101,
      "Hiredate" : "October 3, 1998"
    },
    {
      "name" : "Tom",
      "id" : 1121,
      "Hiredate" : "May 8, 2004"
    }
  ]
}
```

3.2. Index

- It is a collection of different types of documents and their properties.

3.3. Shards

- Indexes are horizontally subdivided into shards.
- This means each shard contains all the properties of document, but contains less number of JSON objects than index.
- The horizontal separation makes shard a fully-functional and independent "index" that can be hosted on any node in the cluster.
- Shards are of two types - Primary and Replica.

3.3.1. Primary Shards

- Primary shard is the original horizontal part of an index and then these primary shards are replicated into replica shards.

3.3.2. Replica Shards

- Elasticsearch allows a user to create replicas of their primary shards.
- Replication not only helps in increasing the availability of data in case of failure, but also improves the performance of searching by carrying out a parallel search operation in these replicas.
- Replicas and their corresponding primary shards are not stored in the same node. This helps in recovery of data in case of node failure.

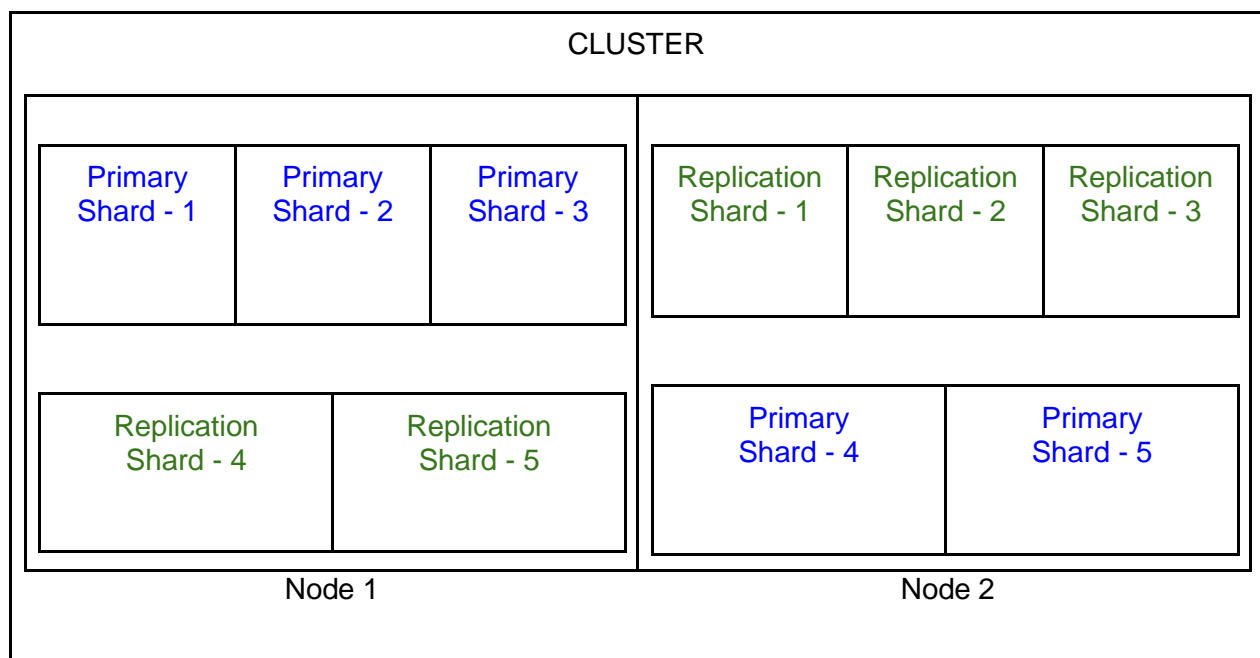
3.4. Nodes

- It refers to a single running instance of Elasticsearch.
- A node is a single server that is part of your cluster, stores your data, and participates in the cluster's indexing and search capabilities.
- A node can have multiple primary shards and replica shards.
- By default, each index in Elasticsearch is allocated 5 primary shards and 1 replica which means that if you have at least two nodes in your cluster, your index will have 5 primary shards and another 5 replica shards (1 complete replica) for a total of 10 shards per index.
- if there are no other Elasticsearch nodes currently running on your network, starting a single node will by default form a new single-node cluster named *elasticsearch*.

3.5. Cluster

- It is a collection of one or more nodes.
- Cluster provides collective indexing and search capabilities across all the nodes for entire data.

Fig : Architecture of Elasticsearch



4. Indices API

4.1. Create Index

- The following command creates the index named "customer" using the PUT verb.

```
curl -X PUT "localhost:9200/customer"
```

4.2. Delete Index

- The following command deletes the index named "customer" using the DELETE verb.

```
curl -X DELETE "localhost:9200/customer"
```

4.3. Get Index

- The following command fetched the index details such as mappings, settings and aliases.

```
curl -XGET "http://localhost:9200/customer"
```

4.4. Index exist

- Used to check if the index (indices) exists or not. For example :

```
curl -XHEAD "http://localhost:9200/customer"
```

- The HTTP status code indicates if the index exists or not. A *404* means it does not exist, and *200* means it does.

4.5. Index Settings

- Each index created can have specific settings associated with it, defined in the body.

```
curl -X PUT "localhost:9200/customer" -H 'Content-Type: application/json' -d'
{  "settings" : {
    "index" : {
      "number_of_shards" : 3,
      "number_of_replicas" : 2
    }
  }
}'
```

- Default for number_of_shards is 5.
- Default for number_of_replicas is 1 (ie one replica for each primary shard)

4.5. Get Settings

- The get settings API allows to retrieve settings of index/indices:

```
curl -X GET "localhost:9200/customer/_settings"
```

4.6. Index Mappings

- The PUT mapping API allows you to add fields to an existing index or to change search only settings of existing fields.

```
curl -X PUT "localhost:9200/customer/_mapping/_doc" -H 'Content-Type: application/json' -d'
{
  "properties": {
    "email": {
      "type": "keyword"
    },
    "age": {
      "type": "integer"
    }
  }
}'
```

5. Mapping

5.1. Field Data-types

5.1.1 Numeric Data-types

- *integer*
- *long*
- *short*
- *byte*
- *double*
- *float*
- *half_float*
- *scaled_float*

5.1.2 String Data-types

- *text*
- *keyword*

5.1.3 Range Data-types

- *integer_range*
- *float_range*
- *long_range*
- *double_range*
- *date_range*

5.1.3 Other Data-types

- *date*
- *boolean*
- *binary*
- Array : Array support does not require a dedicated type.
- *object* : for single JSON object.
- *nested* : for arrays of JSON objects
- *geo_point* : for lat/lon points
- *geo_shape* : for complex shapes like polygons
- *ip* : for IPv4 and IPv6 addresses
- *completion* : to provide auto-complete suggestions
- *token_count*
- *murmur3*
- *mapper-murmur3*
- *join*
- *alias*
- *percolator*

5.2. Dynamic Mapping

- Elasticsearch provides a user-friendly mechanism for the automatic creation of mapping.
- A user can post the data directly to any undefined mapping and Elasticsearch will automatically create the mapping, which is called dynamic mapping.
- Format :

```
PUT index_name/_doc/_id
{
  "field_name": "value"
}
```

- eg :

```
PUT data/_doc/1
{
  "firstname": "Peter",
  "lastname": "Parkar",
  "age": 21,
  "encryption_key": "ecaescEWDEecwWEreW531",
  "experience": {
    "gte": 3,
    "lte": 5
  },
  "hire_date": "2018-08-29 12:20:57",
  "ip_address": "127.0.0.0",
  "isRegular": true,
  "location": [
    42.5687,
    84.4511
  ],
  "manager": {
    "age": 31,
    "name": {
      "first": "Tony",
      "last": "Stark"
    }
  },
  "position": "Trainee",
  "salary": 15000.5,
  "tags": [
    "Genius",
    "Lazy"
  ]
}
```

6. Query DSL (Domain Specific Language)

- In Elasticsearch, searching is carried out by using query based on JSON.

6.1. Match all query

- matches all documents, giving them all a `_score` of 1.0.
- To fetch all documents in all indices:

```
GET /_search
{
  "query": {
    "match_all": {}
  }
}
```

- To fetch all documents in particular index:

```
GET index_name/_search
{
  "query": {
    "match_all": {}
  }
}
```

6.2. Match query

- matches a text or phrase with the values of one field.
- Eg:

```
GET /_search
{
  "query": {
    "match": {
      "city": "pune"
    }
  }
}
```

6.3. Multi Match query

- matches a text or phrase with the values of one or more fields.
- Eg:

```
GET /_search
{
  "query":{
    "multi_match" : {
      "query": "hyderabad",
      "fields": [ "city", "state" ]
    }
  }
}
```

6.4. Match None query

- matches no documents.
- Eg:

```
GET /_search
{
  "query":{
    "match_none" : { }
  }
}
```

6.5. Match Phrase query

- analyzes the text and creates a phrase query out of the analyzed text.
- Eg:

```
GET /_search
{
  "query": {
    "match_phrase" : {
      "message" : "this is a test"
    }
  }
}
```

6.6. Term query (Deprecated in 7.3.0)

- The term query finds documents that contain the exact term (without analyzing), specified in the inverted index.
- Eg:

```
POST _search
{
  "query": {
    "term" : { "user" : "Kimchy" }
  }
}
```

6.7. Range query

- Matches documents with fields that have terms within a certain range..
- Eg:

```
GET _search
{
  "query": {
    "range" : {
      "age" : {
        "gte" : 10,
        "lte" : 20
      }
    }
  }
}
```

6.8. Fuzzy query

- Autocomplete and autocorrect facility.
- Eg:

```
GET /_search
{
  "query": {
    "fuzzy" : { "user" : "ki" }
  }
}
```

6.9. Wildcard query

- Matches documents that have fields matching a wildcard expression.
- Supported wildcards are :
 - * : matches any character sequence (including the empty one),
 - ? : matches any single character.
- Eg:

```
GET /_search
{
  "query": {
    "wildcard" : { "user" : "ki*y" }
  }
}
```

6.10. Bool query

- matches documents using boolean combinations.
 - And : must
 - Or : should
 - Not : must_not
- Eg:

```
GET /bank/_search
{
  "query": {
    "bool": {
      "must": [
        { "match": { "address": "mill" } },
        { "match": { "address": "lane" } }
      ]
    }
  }
}
```

```
GET /bank/_search
{
  "query": {
    "bool": {
      "should": [
        { "match": { "address": "mill" } },
        { "match": { "address": "lane" } }
      ]
    }
  }
}
```

6.11. Nested query

- allows to query nested documents .
- Eg:

```
PUT /my_index
{
  "mappings": {
    "type1" : {
      "properties" : {
        "obj1" : {
          "type" : "nested"
        }
      }
    }
  }
}
```

6.12. Has Child query

- The has_child filter accepts a query and the child type to run against, and results in parent documents that have child docs matching the query
- Eg:

```
GET /_search
{
  "query": {
    "has_child" : {
      "type" : "blog_tag",
      "query" : {
        "term" : {
          "tag" : "something"
        }
      }
    }
  }
}
```

6.13. Has Parent query

- This query returns child documents which associated parents have matched.
- Eg:

```
GET /_search
{
  "query": {
    "has_parent" : {
      "parent_type" : "blog",
      "query" : {
        "term" : {
          "tag" : "something"
        }
      }
    }
  }
}
```

6.14. Geo Shape query

- Filter documents indexed using the geo_shape type.
- Eg:

```
GET /example/_search
{
  "query": {
    "bool": {
      "must": {
        "match_all": {}
      },
      "filter": {
        "geo_shape": {
          "location": {
            "shape": {
              "type": "envelope",
              "coordinates" : [[13.0, 53.0], [14.0, 52.0]]
            },
            "relation": "within"
          }
        }
      }
    }
  }
}
```

6.15. Geo Bounding box query

- A query allowing to filter hits based on a point location using a bounding box.
- Eg:

```
GET /_search
{
  "query": {
    "bool" : {
      "must" : {
        "match_all" : {}
      },
      "filter" : {
        "geo_bounding_box" : {
          "pin.location" : {
            "top_left" : {
              "lat" : 40.73,
              "lon" : -74.1
            },
            "bottom_right" : {
              "lat" : 40.01,
              "lon" : -71.12
            }
          }
        }
      }
    }
  }
}
```

6.16. Geo Distance query

- These queries help to find out schools or any other geographical object near to any location.
- Eg:

```
GET /_search
{
  "query":{
    "filtered":{
      "filter":{
        "geo_distance":{ "distance":"100km",
          "location":[32.052098, 76.649294]
        }
      }
    }
  }
}
```

6.17. Geo Polygon query

- A query allowing to include hits that only fall within a polygon of points.

- Eg:

```
GET /_search
{
  "query": {
    "bool": {
      "must": {
        "match_all": {}
      },
      "filter": {
        "geo_polygon": {
          "person.location": {
            "points": [
              {"lat": 40, "lon": -70},
              {"lat": 30, "lon": -80},
              {"lat": 20, "lon": -90}
            ]
          }
        }
      }
    }
  }
}
```

6.18. Query String query

- A query that uses a query parser in order to parse its content.
- Eg:

```
GET /_search
{
  "query": {
    "query_string": {
      "default_field": "content",
      "query": "this AND that OR thus"
    }
  }
}
```

- It supports field name prefixes, wildcard characters, or other "advanced" features.

7. Aggregations

- The aggregations framework helps provide aggregated data based on a search query.

7.1. Metric Aggregations

- Avg Aggregation

```
POST /exams/_search?size=0
{
  "aggs": {
    "avg_grade": {
      "avg": {
        "field": "grade"
      }
    }
  }
}
```

- Weighted Avg Aggregation

```
POST /exams/_search
{
  "size": 0,
  "aggs": {
    "weighted_grade": {
      "weighted_avg": {
        "value": {
          "field": "grade"
        },
        "weight": {
          "field": "weight"
        }
      }
    }
  }
}
```

- Cardinality Aggregation

```
POST /sales/_search?size=0
{
  "aggs" : {
    "type_count" : {
      "cardinality" : {
        "field" : "type"
      }
    }
  }
}
```

- Extended Stats Aggregation

```
GET /exams/_search
{
  "size": 0,
  "aggs" : {
    "grades_stats" : { "extended_stats" : { "field" : "grade" } }
  }
}
```

- Geo Bounds Aggregation

```
POST /museums/_search?size=0
{
  "query" : {
    "match" : { "name" : "musée" }
  },
  "aggs" : {
    "viewport" : {
      "geo_bounds" : {
        "field" : "location",
        "wrap_longitude" : true
      }
    }
  }
}
```

- Geo Centroid Aggregation

```
PUT /museums
{
  "mappings": {
    "doc": {
      "properties": {
        "location": {
          "type": "geo_point"
        }
      }
    }
  }
}
```

- Max Aggregation

```
POST /sales/_search?size=0
{
  "aggs" : {
    "max_price" : {
      "max" : {
        "field" : "price"
      }
    }
  }
}
```

- Min Aggregation

```
POST /sales/_search?size=0
{
  "aggs" : {
    "max_price" : {
      "max" : {
        "field" : "price"
      }
    }
  }
}
```

- Percentiles Aggregation

```
GET latency/_search
{
  "size": 0,
  "aggs": {
    "load_time_outlier": {
      "percentiles": {
        "field": "load_time"
      }
    }
  }
}
```

- Percentile Ranks Aggregation

```
GET latency/_search
{
  "size": 0,
  "aggs": {
    "load_time_ranks": {
      "percentile_ranks": {
        "field": "load_time",
        "values": [500, 600]
      }
    }
  }
}
```

- Stats Aggregation

```
POST /exams/_search?size=0
{
  "aggs": {
    "grades_stats": {
      "stats": {
        "field": "grade"
      }
    }
  }
}
```

- Sum Aggregation

```
POST /sales/_search?size=0
{
  "query" : {
    "constant_score" : {
      "filter" : {
        "match" : { "type" : "hat" }
      }
    }
  },
  "aggs" : {
    "hat_prices" : { "sum" : { "field" : "price" } }
  }
}
```

- Value Count Aggregation

```
POST /sales/_search?size=0
{
  "aggs" : {
    "types_count" : { "value_count" : { "field" : "type" } }
  }
}
```

8. Index Aliases API

- The index aliases API allows aliasing an index with a name, with all APIs automatically converting the alias name to the actual index name.

- *Eg. Associating the alias alias1 with index test1:*

```
POST /_aliases
{
  "actions" : [
    { "add" : {
      "index" : "test1",
      "alias" : "alias1"
    }
  ]
}
```

- *Eg. Removing the alias alias1 with index test1:*

```
POST /_aliases
{
  "actions" : [
    { "remove" : {
      "index" : "test1",
      "alias" : "alias1"
    }
  ]
}
```

- *Eg. Rename an alias by doing 'remove' then 'add' operation within the same API :*

```
POST /_aliases
{
  "actions" : [
    { "remove" : {
      "index" : "test1", "alias" : "alias1" }
    },
    { "add" : {
      "index" : "test2", "alias" : "alias1" }
    }
  ]
}
```

- *Eg. Associating an alias to more than one index :*

```

POST                                                                    /_aliases
{
  "actions"                                                                :
  {
    "add"      : { "index" : "test1", "alias" : "alias1" },
    "add"      : { "index" : "test2", "alias" : "alias1" }
  }
}

```

OR

```

POST                                                                    /_aliases
{
  "actions"                                                                :
  {
    "add" : { "indices" : ["test1", "test2"], "alias" : "alias1" }
  }
}

```

- Eg. Associating an alias to more than one index that shares a common name:

```

POST                                                                    /_aliases
{
  "actions"                                                                :
  {
    "add" : { "index" : "test*", "alias" : "all_test_indices" }
  }
}

```

- Eg. Swap an index/ rename Index:

```

PUT                                                                    test_2
POST                                                                    /_aliases
{
  "actions"                                                                :
  {
    "add": { "index": "test_2", "alias": "test" },
    "remove_index": { "index": "test" }
  }
}

```